

DATA EXPLORATION & PREPROCESSING A data exploration and preprocessing notebook or report that analyzes the dataset, handles missing values, and prepares the data for modeling

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style='whitegrid')
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (accuracy_score, precision_score,
recall_score, f1_score,
                           roc_auc_score, classification_report,
                           confusion_matrix)
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
import joblib

from imblearn.over_sampling import SMOTE

# 1.1 Load data & quick overview

df = pd.read_csv(r"C:\Users\Pravallika\Downloads\Customerdata.csv")
print("Shape:", df.shape)
display(df.head())
display(df.info())
display(df.describe(include='all').T)

# Target check (example: 'Churn' column with Yes/No)
print(df['Churn'].value_counts(dropna=False))
sns.countplot(x='Churn', data=df)
plt.title('Churn Distribution'); plt.show()

Shape: (7043, 21)

   customerID  gender  SeniorCitizen Partner Dependents  tenure
PhoneService \
0    7590-VHVEG  Female                  0      Yes        No       1
No
1    5575-GNVDE    Male                  0      No        No      34
Yes
2    3668-QPYBK    Male                  0      No        No       2
Yes
3    7795-CFOCW    Male                  0      No        No      45
```

No						
4	9237-HQITU	Female		0	No	No
Yes						2
	MultipleLines	InternetService	OnlineSecurity	...		
DeviceProtection	\					
0	No phone service		DSL		No	...
No						
1		No	DSL		Yes	...
Yes						
2		No	DSL		Yes	...
No						
3	No phone service		DSL		Yes	...
Yes						
4		No	Fiber optic		No	...
No						
	TechSupport	StreamingTV	StreamingMovies		Contract	
PaperlessBilling	\					
0	No	No		No	Month-to-month	
Yes						
1	No	No		No	One year	
No						
2	No	No		No	Month-to-month	
Yes						
3	Yes	No		No	One year	
No						
4	No	No		No	Month-to-month	
Yes						
	PaymentMethod	MonthlyCharges	TotalCharges		Churn	
0	Electronic check	29.85	29.85		No	
1	Mailed check	56.95	1889.50		No	
2	Mailed check	53.85	108.15		Yes	
3	Bank transfer (automatic)	42.30	1840.75		No	
4	Electronic check	70.70	151.65		Yes	

[5 rows x 21 columns]

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 5   tenure          7043 non-null   int64 
```

```

6 PhoneService      7043 non-null   object
7 MultipleLines     7043 non-null   object
8 InternetService   7043 non-null   object
9 OnlineSecurity    7043 non-null   object
10 OnlineBackup      7043 non-null   object
11 DeviceProtection  7043 non-null   object
12 TechSupport       7043 non-null   object
13 StreamingTV        7043 non-null   object
14 StreamingMovies    7043 non-null   object
15 Contract          7043 non-null   object
16 PaperlessBilling  7043 non-null   object
17 PaymentMethod      7043 non-null   object
18 MonthlyCharges    7043 non-null   float64
19 TotalCharges      7032 non-null   float64
20 Churn              7043 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB

```

None


		count	unique	top	freq	
mean \	customerID	7043	7043	3186-AJIEK	1	NaN
	gender	7043	2	Male	3555	NaN
	SeniorCitizen	7043.0	NaN	NaN	NaN	0.162147
	Partner	7043	2	No	3641	NaN
	Dependents	7043	2	No	4933	NaN
	tenure	7043.0	NaN	NaN	NaN	32.371149
	PhoneService	7043	2	Yes	6361	NaN
	MultipleLines	7043	3	No	3390	NaN
	InternetService	7043	3	Fiber optic	3096	NaN
	OnlineSecurity	7043	3	No	3498	NaN
	OnlineBackup	7043	3	No	3088	NaN
	DeviceProtection	7043	3	No	3095	NaN
	TechSupport	7043	3	No	3473	NaN
	StreamingTV	7043	3	No	2810	NaN
	StreamingMovies	7043	3	No	2785	NaN

Contract	7043	3	Month-to-month	3875		NaN
PaperlessBilling	7043	2		Yes	4171	NaN
PaymentMethod	7043	4	Electronic check	2365		NaN
MonthlyCharges	7043.0	NaN		NaN	NaN	64.761692
TotalCharges	7032.0	NaN		NaN	NaN	2283.300441
Churn	7043	2		No	5174	NaN
		std	min	25%	50%	75%
max						
customerID		NaN	NaN	NaN	NaN	NaN
NaN						
gender		NaN	NaN	NaN	NaN	NaN
NaN						
SeniorCitizen	0.368612	0.0	0.0	0.0	0.0	0.0
1.0						
Partner		NaN	NaN	NaN	NaN	NaN
NaN						
Dependents		NaN	NaN	NaN	NaN	NaN
NaN						
tenure	24.559481	0.0	9.0	29.0	55.0	
72.0						
PhoneService		NaN	NaN	NaN	NaN	NaN
NaN						
MultipleLines		NaN	NaN	NaN	NaN	NaN
NaN						
InternetService		NaN	NaN	NaN	NaN	NaN
NaN						
OnlineSecurity		NaN	NaN	NaN	NaN	NaN
NaN						
OnlineBackup		NaN	NaN	NaN	NaN	NaN
NaN						
DeviceProtection		NaN	NaN	NaN	NaN	NaN
NaN						
TechSupport		NaN	NaN	NaN	NaN	NaN
NaN						
StreamingTV		NaN	NaN	NaN	NaN	NaN
NaN						
StreamingMovies		NaN	NaN	NaN	NaN	NaN
NaN						
Contract		NaN	NaN	NaN	NaN	NaN
NaN						
PaperlessBilling		NaN	NaN	NaN	NaN	NaN
NaN						

```

PaymentMethod          NaN      NaN      NaN      NaN      NaN
NaN
MonthlyCharges        30.090047  18.25    35.5    70.35    89.85
118.75
TotalCharges         2266.771362  18.8    401.45   1397.475  3794.7375
8684.8
Churn                NaN      NaN      NaN      NaN      NaN
NaN
Churn
No       5174
Yes      1869
Name: count, dtype: int64

```



```

# Missing values
missing = df.isnull().sum().sort_values(ascending=False)
print("Missing values per column:\n", missing[missing>0])

# Visual missingness
plt.figure(figsize=(10,4))
sns.heatmap(df.isnull(), cbar=False)
plt.title('Missing Data Heatmap'); plt.show()

```


```

# Price/charge-like columns: outlier check example
for col in ['tenure', 'MonthlyCharges', 'TotalCharges']:
    if col in df.columns:
        plt.figure(figsize=(6,2))
        sns.boxplot(x=df[col])
        plt.title(f'Boxplot: {col}')
        plt.show()


# Correlation heatmap for numeric features
plt.figure(figsize=(10,8))
sns.heatmap(df.select_dtypes(include=[np.number]).corr(), annot=True,
            fmt=".2f", cmap='coolwarm')
plt.title('Numeric feature correlations'); plt.show()

Missing values per column:
TotalCharges      11
dtype: int64

```




Boxplot: tenure




Boxplot: MonthlyCharges



Boxplot: TotalCharges





```
# 1.2 Cleaning & Feature Engineering

if 'TotalCharges' in df.columns and df['TotalCharges'].dtype == 'object':
    df['TotalCharges'] = pd.to_numeric(df['TotalCharges'],
errors='coerce')

# Impute missing values
num_cols =
df.select_dtypes(include=['int64','float64']).columns.tolist()
cat_cols =
df.select_dtypes(include=['object','category']).columns.tolist()
# remove target from cat_cols
if 'Churn' in cat_cols:
    cat_cols.remove('Churn')
```

```

num_imputer = SimpleImputer(strategy='median')
cat_imputer = SimpleImputer(strategy='most_frequent')

df[num_cols] = pd.DataFrame(num_imputer.fit_transform(df[num_cols]),
columns=num_cols)
df[cat_cols] = pd.DataFrame(cat_imputer.fit_transform(df[cat_cols]),
columns=cat_cols)

# Feature engineering examples:
if 'tenure' in df.columns:
    df['tenure_group'] = pd.cut(df['tenure'], bins=[-1,3,12,24,48,120], labels=['0-3','4-12','13-24','25-48','49+'])

# Create a numeric indicator for multi-service usage example
service_cols = [c for c in df.columns if
c.lower().startswith(('internet','online','tech','stream'))]
if service_cols:
    df['num_services'] = df[service_cols].apply(lambda row:
sum(row=='Yes') if row.dtype=='O' else np.nan,
axis=1).fillna(0).astype(int)

for c in cat_cols:
    df[c] = df[c].replace(' ', np.nan).fillna(df[c].mode()[0])

```

A machine learning model capable of predicting customer churn

```

# Prepare X, y
target = 'Churn'
y = df[target].map({'Yes':1,'No':0}) if df[target].dtype=='object'
else df[target]
X = df.drop(columns=[target])

# Identify column types for transformer
numeric_features =
X.select_dtypes(include=[np.number]).columns.tolist()
categorical_features =
X.select_dtypes(include=['object','category']).columns.tolist()

# Preprocessing pipelines
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
```

```

        ('onehot', OneHotEncoder(handle_unknown='ignore',
sparse_output=False))

])

preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_features),
    ('cat', categorical_transformer, categorical_features)
], remainder='drop')

#Train/validation/test split

X_trainval, X_test, y_trainval, y_test = train_test_split(X, y,
test_size=0.20, random_state=42, stratify=y)
X_train, X_val, y_train, y_val = train_test_split(X_trainval,
y_trainval, test_size=0.1875, random_state=42, stratify=y_trainval)
print("Shapes -> Train:", X_train.shape, "Val:", X_val.shape, "Test:",
X_test.shape)

Shapes -> Train: (4577, 22) Val: (1057, 22) Test: (1409, 22)

# We'll build pipelines that include SMOTE after preprocessing; use imblearn Pipeline
from imblearn.pipeline import Pipeline as ImbPipeline
from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)

# Logistic Regression (interpretable baseline)
pipe_lr = ImbPipeline(steps=[
    ('preprocessor', preprocessor),
    ('smote', smote),
    ('clf', LogisticRegression(max_iter=1000,
class_weight='balanced'))
])
pipe_lr.fit(X_train, y_train)

# Random Forest (powerful, interpretable via importances)
pipe_rf = ImbPipeline(steps=[
    ('preprocessor', preprocessor),
    ('smote', smote),
    ('clf', RandomForestClassifier(random_state=42,
class_weight='balanced', n_jobs=-1))
])
param_grid_rf = {
    'clf__n_estimators': [100, 200],
    'clf__max_depth': [None, 10, 20]
}
grid_rf = GridSearchCV(pipe_rf, param_grid_rf, cv=3, scoring='f1',

```

```

n_jobs=-1)
grid_rf.fit(X_train, y_train)

# XGBoost (often best performance)
pipe_xgb = ImbPipeline(steps=[
    ('preprocessor', preprocessor),
    ('smote', smote),
    ('clf', XGBClassifier(eval_metric='logloss', random_state=42))
])
param_grid_xgb = {
    'clf_n_estimators': [100, 200],
    'clf_max_depth': [3,6]
}
grid_xgb = GridSearchCV(pipe_xgb, param_grid_xgb, cv=3, scoring='f1',
n_jobs=-1)
grid_xgb.fit(X_train, y_train)

print("Best RF params:", grid_rf.best_params_)
print("Best XGB params:", grid_xgb.best_params_)

```

Model Evaluation An evaluation of model performance using appropriate metrics such as accuracy, precision, recall, F1 score, etc.

```

# Utility to evaluate and print metrics
def evaluate_model(name, model, X, y):
    preds = model.predict(X)
    probs = model.predict_proba(X)[:,1] if hasattr(model,
"predict_proba") else None
    acc = accuracy_score(y, preds)
    prec = precision_score(y, preds)
    rec = recall_score(y, preds)
    f1 = f1_score(y, preds)
    roc = roc_auc_score(y, probs) if probs is not None else None
    print(f"--- {name} ---")
    print(f"Accuracy: {acc:.4f} Precision: {prec:.4f} Recall:
{rec:.4f} F1: {f1:.4f} ROC-AUC: {roc:.4f}")
    print(classification_report(y, preds))
    cm = confusion_matrix(y, preds)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'Confusion Matrix - {name}'); plt.xlabel('Pred');
    plt.ylabel('True'); plt.show()
    return
{'accuracy':acc,'precision':prec,'recall':rec,'f1':f1,'roc':roc}

# Evaluate on validation to choose best model
res_lr = evaluate_model("LogisticRegression (val)", pipe_lr, X_val,
y_val)
res_rf = evaluate_model("RandomForest (val)", grid_rf.best_estimator_,


```

```

X_val, y_val)
res_xgb = evaluate_model("XGBoost (val)", grid_xgb.best_estimator_,
X_val, y_val)

--- LogisticRegression (val) ---
Accuracy: 0.7881 Precision: 0.5782 Recall: 0.7393 F1: 0.6489 ROC-
AUC: 0.8556
      precision    recall   f1-score  support
          0         0.90     0.81     0.85      777
          1         0.58     0.74     0.65      280
   accuracy          0.74     0.77     0.75      1057
  macro avg          0.74     0.77     0.75      1057
weighted avg          0.81     0.79     0.80      1057

```




```


--- RandomForest (val) ---
Accuracy: 0.7682 Precision: 0.5434 Recall: 0.7821 F1: 0.6413 ROC-
AUC: 0.8505
      precision    recall   f1-score  support
          0         0.91     0.76     0.83      777

```

	1	0.54	0.78	0.64	280
accuracy				0.77	1057
macro avg		0.73	0.77	0.74	1057
weighted avg		0.81	0.77	0.78	1057



--- XGBoost (val) ---					
	precision	recall	f1-score	support	
0	0.86	0.91	0.88	777	
1	0.69	0.58	0.63	280	
accuracy			0.82	1057	
macro avg	0.77	0.74	0.75	1057	
weighted avg	0.81	0.82	0.81	1057	



```


val_scores = {'LR':res_lr['f1'], 'RF':res_rf['f1'],
'XGB':res_xgb['f1']}
best_key = max(val_scores, key=val_scores.get)
best_model = {'LR':pipe_lr, 'RF':grid_rf.best_estimator_,
'XGB':grid_xgb.best_estimator_}[best_key]
print("Best model selected:", best_key, "with F1:",
val_scores[best_key])

Best model selected: LR with F1: 0.6489028213166145

#Evaluate on test set
res_test = evaluate_model(f"BestModel-{best_key}(test)", best_model,
X_test, y_test)

--- BestModel-LR(test) ---
Accuracy: 0.7722 Precision: 0.5535 Recall: 0.7326 F1: 0.6306 ROC-
AUC: 0.8452
      precision    recall   f1-score   support
      0          0.89     0.79     0.84     1035
  
```

	1	0.55	0.73	0.63	374
accuracy				0.77	1409
macro avg		0.72	0.76	0.73	1409
weighted avg		0.80	0.77	0.78	1409



```
# I
def get_feature_names(column_transformer):
    # numeric features
    num = column_transformer.transformers_[0][2]
    cat = column_transformer.transformers_[1][2]
    # onehot feature names
    ohe =
column_transformer.named_transformers_['cat'].named_steps['onehot']
    cat_names =
ohe.get_feature_names_out(column_transformer.transformers_[1][2])
    return list(num) + list(cat_names)
```


```

if best_key in ['RF', 'XGB']:
    # extract preprocessor from pipeline
    pre = best_model.named_steps['preprocessor']
    feat_names = list(numeric_features) +
list(best_model.named_steps['preprocessor'].named_transformers_['cat'].
    named_steps['onehot'].get_feature_names_out(categorical_features))
    importances = best_model.named_steps['clf'].feature_importances_
    feat_imp = pd.Series(importances,
index=feat_names).sort_values(ascending=False).head(20)
    plt.figure(figsize=(8,6))
    sns.barplot(x=feat_imp.values, y=feat_imp.index)
    plt.title('Top 20 Feature Importances'); plt.show()

import shap

X_background =
best_model.named_steps['preprocessor'].transform(X_train.sample(200))
explainer =
shap.LinearExplainer( model=best_model.named_steps['clf'],masker=X_bac
kground)
X_to_explain = X_background
shap_values = explainer.shap_values(X_to_explain)
feat_names =
best_model.named_steps['preprocessor'].get_feature_names_out()
shap.summary_plot(shap_values, X_to_explain, feature_names=feat_names)

```



```
# Save best model
joblib.dump(grid_rf.best_estimator_, 'best_churn_model.joblib')
# Assuming grid_rf is the best model
```

```
# Predict function for new raw dataframe (must have same columns)
def predict_new(df_new, model_path='best_churn_model.joblib'):
    model = joblib.load(model_path)
    preds = model.predict(df_new)
    # Check if the model has predict_proba before accessing it
    probs = model.predict_proba(df_new)[:,1] if hasattr(model,
'predict_proba') else None
    return preds, probs

# Example usage:
# Assuming X_test is a raw dataframe before preprocessing
new_df = X_test.iloc[[0]] # raw (not preprocessed) row
p, pr = predict_new(new_df)
print('Pred:', p, 'Prob:', pr)

Pred: [0] Prob: [0.18218309]
```